

1 Strategies for Merging Earthquake Catalogs: A Literature Review

The construction of unified earthquake catalogs from multiple heterogeneous source networks is one of the central methodological challenges of observational seismology. The need arises at every level of the discipline: global hazard models require decades of stable, magnitude-homogeneous records spanning multiple agencies (Pagani et al., 2014); regional seismic-rate studies must combine specialist national networks with global bulletins to fill spatial coverage gaps (Woessner and Wiemer, 2005); historical and pre-instrumental periods demand careful integration of macroseismic and early-instrumental data (Grünthal, 1998; Rovida et al., 2016); and the growth of dense, heterogeneous strong-motion and nodal-sensor networks creates new challenges at sub-local scales (Share et al., 2022).

Despite the ubiquity of the problem, no single standard algorithmic pipeline has emerged across the community. Different groups have developed independent solutions—from the rule-based proximity approaches of the Northern California Earthquake Data Center (2024) and the ISC bulletin (Bondár and Storchak, 2011), to the graph-theoretic frameworks of Patton et al. (2016), to the physics-based nearest-neighbour discriminant of Vorobieva et al. (2022)—each optimised for different data volumes, network configurations, and seismicity regimes. The goal of this review is to synthesise the methodological state of the art into a coherent framework that can inform automated merging implementations.

1.1 The Two-Step Decomposition

Any robust merging pipeline must separate two operations that are logically distinct but are frequently conflated in practice:

1. **Event association and deduplication** — determining which records across source catalogs describe the *same physical earthquake* and grouping them into equivalence classes.
2. **Parameter selection and homogenisation** — from each equivalence class, deriving a single best origin time, hypocentre, and magnitude, and expressing all magnitudes on a common scale.

The distinction matters because applying a magnitude conversion formula *before* the association step—for instance, to make thresholds independent of magnitude scale—introduces feedback loops: converted magnitudes may shift an event across an association threshold, altering which pairs are grouped, which in turn changes the sample from which conversion coefficients are estimated (Storchak et al., 2013; Grünthal et al., 2013). The correct order is always: associate first using raw reported values; homogenise second using the confirmed groups.

1.2 Historical Development and the Heterogeneous Network Problem

1.2.1 The pre-digital era: 1900–1970

Early global earthquake catalogs—the Gutenberg-Richter Seismicity of the Earth (Gutenberg and Richter, 1954), the Abe catalog of large historical earthquakes (Abe, 1981), and the pre-1964 portion of the Engdahl-van der Hilst-Buland (EHB) dataset (Engdahl et al., 1998)—were assembled by manual examination of seismogram archives, station bulletins, and published macroseismic accounts. Cross-network event identification was performed by inspection, using timing differences of several minutes and epicentral discrepancies of hundreds of kilometres as acceptable tolerances for the analogue recording era (Bondár and Storchak, 2011).

The International Seismological Summary (ISS, predecessor to the ISC bulletin) began assembling global phase data in 1918. Its practice of retaining all contributing agency solutions alongside a centralised re-location—still followed by the ISC today—established the principle of *solution provenance* that underpins modern best practice (Storchak et al., 2020).

1.2.2 The WWSSN and NEIC era: 1960–1990

The deployment of the World-Wide Standardized Seismograph Network (WWSSN) from 1961 onward provided the first globally uniform recording infrastructure and enabled systematic relocation of events at $M \geq 5$ using common velocity models (Engdahl et al., 1998). The 1964 establishment of the National Earthquake Information Center (NEIC, initially the National Earthquake Information Service) created the first dedicated fusion centre for global earthquake data, introducing the concept of a *preferred solution* from a ranked list of contributing agencies (USGS, 2025).

The ANSS Composite catalog, which consolidates the NEIC global bulletin with authoritative regional networks across North America, employs a proximity rule of 16 s in time and 100 km in distance to identify duplicate reports submitted by different networks for the same event (Northern California Earthquake Data Center, 2024). This rule was formulated for the digital WWSSN era and remains the most precisely documented operational threshold in any publicly available system documentation.

1.2.3 The digital broadband era: 1990–present

The transition to digital broadband recording in the late 1980s and 1990s, the expansion of the IRIS/IDA Global Seismographic Network (GSN), and the growth of the FDSN data-exchange standard (ETH Zürich Seismology / QuakeML Working Group, 2015) transformed the scale of the deduplication problem. By 2010, the ISC bulletin was receiving origin reports from more than 130 contributing agencies worldwide, each applying different velocity models, phase identification algorithms, and location programs (Storchak et al., 2020). The ISC-GEM Global Instrumental Earthquake Catalog (Storchak et al., 2013; Di Giacomo et al., 2018), released in 2013 and subsequently extended through 2021, represents the most comprehensive attempt to retrospectively homogenise this heterogeneous archive for events $M \geq 5.5$.

The proliferation of dense regional networks—GeoNet in Aotearoa New Zealand (GeoNet, 2025a), Hi-net in Japan (Obara et al., 2005), the Southern California Seismic Network (Hutton et al., 2010), and the European EPOS infrastructure (EPOS, 2025)—has simultaneously improved detection thresholds for small events ($M < 3$) while creating new challenges for merging: specialist local networks now routinely detect thousands of microseismic events per year that global agencies never report, making one-to-one association across scales difficult (Woessner and Wiemer, 2005; Illsley-Kemp and Mestel, 2025).

1.3 Event Association: Spatio-Temporal Matching Algorithms

1.3.1 Physical basis for proximity thresholds

The rationale for any spatio-temporal matching window is that the reported origin time and hypocentre of two records for the same physical earthquake will differ only by the combined effect of (i) timing errors at each recording network, (ii) location uncertainties arising from network geometry and velocity model errors, and (iii) any deliberate choices such as depth fixing.

For digital networks with GPS-synchronised clocks, origin time errors are typically well below 1 s (Bondár and Storchak, 2011). Location errors depend strongly on network geometry: the number of stations, the azimuthal gap (the largest angular gap between contributing stations), the secondary azimuthal gap, and the distance to the nearest station are the dominant predictors. Bondár et al. (2004) expressed this as network criteria under which an epicentre is accurate to a stated level, rather than as a formula for the error. For local networks, an event recorded by at least ten stations within 250 km, with an azimuthal gap below 110° , a secondary azimuthal gap below 160° and at least one station within 30 km of the epicentre is located to within 5 km at 95% confidence (GT5_{95%}); regional and teleseismic network geometries support only coarser accuracy classes, of tens of kilometres. Events outside these criteria, and above all those outside the network with gaps above 180° , are routinely mislocated by tens of kilometres (Engdahl et al., 1998; Bondár et al., 2004), which is what justifies wider matching windows for them.

Pre-digital era timing errors of several seconds per station translate to origin time uncertainties of 1–10 s for well-recorded events and up to 30–60 s for poorly-recorded events in the pre-1960 period (Bondár and Storchak, 2011). This is the physical basis for the wider windows in table 1.

1.3.2 Window-based forward-sweep

The oldest and most widely deployed method sorts all incoming origin reports by origin time and scans forward with a sliding window. If two reports from *different* networks fall within a fixed time window $\Delta t_{\max}$ and an epicentral distance $\Delta d_{\max}$, they are flagged as candidate duplicates (Northern California Earthquake Data Center, 2024; Woessner and Wiemer, 2005). Table 1 summarises documented threshold choices from major operational systems.

Table 1: Spatio-temporal matching thresholds documented in major operational earthquake catalog systems. Δt = origin-time window, Δd = epicentral distance window. ANSS thresholds apply to inter-network duplicate pairs only; intra-network deduplication is assumed to have been handled upstream by each contributing agency.

System / Study	Δt (s)	Δd (km)	Reference
ANSS Composite catalog	16	100	Northern California Earthquake Data Center (2024)
Typical global build ($M \geq 4.5$)	60	100	Storchak et al. (2013)
Local events ($M < 4$)	30	30	Woessner and Wiemer (2005)
Pre-digital era (pre-1960)	300	200	Bondár and Storchak (2011)
Dead Sea Transform region	60	100	Grigoratos et al. (2023)
Mexico unified catalog	60	≈ 111 (1°)	Sawires et al. (2019)
Turkey homogeneous catalog	60	100	Tan (2021)
Italy CPTI15 pipeline	30–120	50–150	Rovida et al. (2016)
ISC-GEM global ($M \geq 5.5$)	60	100	Di Giacomo et al. (2018)

The 16 s / 100 km rule of the ANSS is the most precisely documented publicly available threshold for a production system. It was chosen empirically to balance false-positive (erroneously merging distinct earthquakes) against false-negative (missing genuine duplicate pairs) rates for the population of events in the North American digital-era database (Northern California Earthquake Data Center, 2024). For teleseismic events the 60 s / 100 km window is widely cited; the ISC uses a comparable criterion internally when auto-grouping agency solutions prior to relocation (Storchak et al., 2020).

Algorithmic complexity and spatial indexing. A naive implementation of the forward sweep compares all ordered pairs (E_i, E_j) with $j > i$, giving $O(n^2)$ complexity for n events. For large catalogs ($n > 10^5$), this is prohibitive. The standard optimisation builds a *spatial index* over event locations: a uniform grid (hash grid), k-d tree, or R-tree, allowing events within $\Delta d_{\max}$ to be retrieved in $O(\log n)$ time. Combined with a time-sorted event list and an early-exit rule when the candidate’s origin time exceeds $t_i + \Delta t_{\max}$, the algorithm achieves $O(n \log n)$ overall complexity in the typical case (Patton et al., 2016).

Grid-based indexes use a fixed cell size calibrated to the distance threshold. At high latitudes, degrees of longitude subtend shorter distances, so the cell size must be adjusted by $\cos(\phi)$ where ϕ is the latitude (Tan, 2021). Near the International

Date Line (longitude $\pm 180^\circ$), grid cell keys must wrap correctly; failure to implement wrapping silently loses candidates for events in the Pacific (Grigoratos et al., 2023).

Adaptive thresholds. Fixed windows perform poorly across the full magnitude range because larger events are subject to greater timing and location uncertainty at distant networks. A common remedy is to scale the effective window with event magnitude and depth:

$$\Delta t_{\text{eff}} = \Delta t_{\text{config}} \cdot f_t(M), \quad \Delta d_{\text{eff}} = \Delta d_{\text{config}} \cdot f_d(M) \cdot g(h), \quad (1)$$

where the multiplier functions below are the heuristic choice of the Catalogue of Catalogues prototype, set so that the windows grow with the location and timing scatter expected at each size and depth; they are not a published ISC-GEM table (ISC-GEM applies fixed windows with expert review, table 1):

$$f_t(M) = \begin{cases} 1.0 & M < 4.0 \\ 1.5 & 4.0 \leq M < 5.5 \\ 2.0 & 5.5 \leq M < 7.0 \\ 3.0 & M \geq 7.0 \end{cases} \quad (2)$$

$$f_d(M) = \begin{cases} 1.0 & M < 4.0 \\ 1.5 & 4.0 \leq M < 5.5 \\ 2.5 & 5.5 \leq M < 7.0 \\ 4.0 & M \geq 7.0 \end{cases} \quad (3)$$

$$g(h) = \begin{cases} 1.0 & h < 100 \text{ km} \\ 1.2 & 100 \leq h < 300 \text{ km} \\ 1.5 & h \geq 300 \text{ km} \end{cases} \quad (4)$$

The depth multiplier $g(h)$ reflects the well-established result that intermediate and deep-focus earthquakes are harder to locate precisely because depth-sensitive phases (pP, sP) are often absent from or misidentified in network phase datasets (Engdahl et al., 1998, 2020). For $M \geq 7$ deep events, the combined multiplier can reach $4.0 \times 1.5 = 6.0$, requiring the spatial search radius to extend to six times the configured baseline threshold.

Early termination and threshold consistency. A subtle but important correctness issue arises when combining the forward-sweep time-sorted early exit with adaptive thresholds. The early-exit condition

$$t_j - t_i > \Delta t_{\text{config}} \cdot f_t(M_i) \quad (5)$$

uses only the anchor event’s magnitude M_i . However, the adaptive matching function (1) uses the *average* of both event magnitudes. If the candidate j has a substantially higher magnitude than i , the correct effective threshold (evaluated at the average magnitude) may exceed the early-exit threshold (evaluated at M_i alone), causing valid matches to be discarded. The conservative fix is to replace $f_t(M_i)$ in

(5) with $f_t(\max(M_i, M_j))$, or equivalently to use the maximum expected candidate magnitude (e.g. $M_i + 2$) as a conservative upper bound when evaluating the exit condition (Patton et al., 2016).

Limitation: dense aftershock sequences. During intense aftershock sequences the inter-event time between consecutive legitimate aftershocks can fall well below the duplicate-detection window. Vorobieva et al. (2022) demonstrated that in the 24 hours following the 2011 Tōhoku M_w 9.0, more than 599 genuine aftershock pairs had inter-event times under 60 s — the same interval used as a duplicate window in the ISC-GEM pipeline. Fixed-window methods therefore cannot discriminate a duplicate pair from an aftershock pair in high-seismicity periods without additional physical information, a problem that grows more severe as network density increases and detection completeness lowers the magnitude threshold into the range where aftershock rates are high (Hainzl et al., 2008).

1.3.3 Nearest-neighbour proximity function

Vorobieva et al. (2022) introduced a physics-based proximity function that discriminates duplicate pairs from aftershock pairs by exploiting their different topological signatures in event-pair space. The method builds on the nearest-neighbour statistics introduced by Zaliapin et al. (2008) for aftershock sequence identification. Genuine duplicates form *isolated pairs* in rescaled event space — their nearest-neighbour distances cluster tightly below a network-error-derived threshold — whereas aftershocks form *branching trees* following Omori-law temporal decay (Omori, 1894; Utsu, 1961).

The proximity between two events i and j (with j following i in time) is defined as

$$\eta(i, j) = t_{ij}^{d_f} r_{ij}^{d_f} 10^{-bm_i}, \quad (6)$$

where t_{ij} is the inter-event time (s), r_{ij} is the epicentral distance (km), $d_f \approx 1.6$ is the fractal dimension of the epicentre distribution (Kagan, 1991), $b \approx 1.0$ is the Gutenberg-Richter b -value (Aki, 1965), and m_i is the magnitude of the preceding event. In this rescaled space, aftershocks cluster near their parent event (small η), while duplicates cluster near zero regardless of the parent-event magnitude (Zaliapin et al., 2008; Vorobieva et al., 2022).

The duplicate-detection threshold in η -space is calibrated against the distribution of location errors for the specific network under analysis, not against earthquake physics. This is the key advantage: the threshold adapts to network capability rather than assuming a fixed proximity tolerance (Vorobieva et al., 2022). Applied to the Tōhoku sequence, the method achieved duplicate identification reliability $> 97\%$ even during periods with aftershock rates of several hundred events per hour.

Parameter estimation. The fractal dimension d_f can be estimated from the slope of the correlation integral of hypocentre coordinates in log-log space (Kagan, 1991). The b -value is estimated via maximum likelihood from the observed frequency-magnitude distribution (Aki, 1965), with magnitude of completeness M_c determined by the Wiemer and Wyss (2000) goodness-of-fit method before fitting.

Both estimates should be computed from the target region’s seismicity, not borrowed from global averages, as both parameters show significant regional variability (Woessner and Wiemer, 2005).

Limitations. The nearest-neighbour method performs less well for very sparse catalogs where the nearest-neighbour distance distribution is dominated by background seismicity rather than aftershock clusters, reducing the bimodal separation that enables discrimination (Hainzl et al., 2008). It also requires a minimum catalog length (typically several years) to estimate b and d_f reliably.

1.3.4 Graph-based connected-components association

The methodologically strongest general-purpose approach constructs an undirected association graph $G = (V, E)$ in which each node $v \in V$ is an origin report and each edge $e \in E$ connects a pair of reports that satisfy a proximity criterion (Patton et al., 2016; Zaliapin and Ben-Zion, 2013). Connected components of G , recoverable in $O(n \alpha(n))$ time via the Union-Find (Disjoint-Set Union) algorithm (Tarjan, 1975) — where α is the inverse Ackermann function and is effectively constant for all practical n — correspond to clusters of reports describing the same physical earthquake.

Algorithm.

1. Sort all origin reports by origin time; build a spatial index (uniform grid, k-d tree, or R-tree) over event hypocentres.
2. For each report E_i , query the index for all reports within $\Delta t_{\max}$ and $\Delta d_{\max}$ from networks *other than* the network of E_i .
3. For each qualifying candidate pair (E_i, E_j) , add edge e_{ij} to G .
4. Find all connected components of G using Union-Find.
5. For each component, apply parameter-selection rules (section 1.9) to derive a single merged event record.

Advantage over the greedy forward sweep. The greedy sweep implicitly assumes that pairwise matching is transitive: if event A matches B, and B matches C, all three are merged. In the sweep implementation this is only guaranteed when events are processed strictly in time order and no candidate is marked processed before its later matches are evaluated. When the sweep marks B as processed upon matching with A, a subsequent evaluation of C against B is never performed, and the A-B-C cluster is incorrectly split (Patton et al., 2016). The graph approach has no such pathology: transitivity is guaranteed by the connected-component structure.

Integer linear programming refinement. Patton et al. (2016) extended the graph approach by applying Integer Linear Programming (ILP) to resolve ambiguous components — cases where a connected component may span two closely spaced but genuinely distinct earthquakes. The ILP minimises the total number of source events k required to explain all edges in a component, subject to the constraint that no two reports assigned to the same source event originate from the same network:

$$\begin{aligned}
& \text{minimise} && k \\
& \text{subject to} && \forall (i, j) \in E : \text{source}(i) = \text{source}(j), \\
& && \forall i, j : \text{net}(i) = \text{net}(j) \Rightarrow \text{source}(i) \neq \text{source}(j).
\end{aligned} \tag{7}$$

Applied to the Chilean catalogue 2010–2017, this two-stage approach detected 817 000+ earthquakes at $> 95\%$ accuracy against a synthetic validation dataset generated by forward modelling known network detection probabilities.

The ILP constraint that *no two reports from the same network belong to the same source event* is the algorithmic expression of what we term the *same-agency duplicate flag*: if two records from the same agency fall within the proximity window, they almost certainly describe two distinct physical earthquakes, not two observations of one (Patton et al., 2016; Vorobieva et al., 2022).

The USGS NEIC Hydra production system uses a graph-association framework conceptually equivalent to the first stage of Patton et al. (2016) (USGS, 2025).

1.4 Location Uncertainty and Its Effect on Matching

1.4.1 Sources of location uncertainty

Hypocentre location uncertainty arises from five primary sources (Bondár et al., 2004; Engdahl et al., 1998):

1. **Phase picking errors.** Manual or automated picks of P , S , and depth phases carry random timing errors (typically 0.1–1 s for digital data, 1–5 s for analogue).
2. **Network geometry.** The azimuthal gap and minimum station distance control the conditioning of the location inversion. A gap $> 180^\circ$ leaves the hypocentre poorly constrained in the ungapped direction; a large minimum distance makes depth poorly constrained.
3. **Velocity model errors.** All location algorithms use an approximation to Earth structure; residual travel-time prediction errors are correlated among nearby stations, biasing the location estimate. Bondár and Storchak (2011) introduced a correlated error covariance model to account for this.
4. **Depth-phase absence.** Without pP , sP , or pwP phases, depth is often fixed at a default value (10, 33, or 100 km depending on agency convention), introducing a systematic depth error that can exceed 50 km (Engdahl et al., 1998, 2020).
5. **Reporting delays and partial phase sets.** Rapid preliminary solutions from real-time networks (e.g. NEIC PDL) use fewer phases than the reviewed bulletin solution and therefore carry larger uncertainties (USGS, 2025).

1.4.2 Quantifying location quality for merging decisions

The ISC bulletin quality scoring system (Bondár and Storchak, 2011) provides a practical framework for weighting competing hypocentre solutions within a merged group. The primary quality indicators are:

- Azimuthal gap: $GAP < 180^\circ$ for acceptable, $< 120^\circ$ for good.
- Secondary azimuthal gap (gap ignoring the single most distant station): $< 180^\circ$ is the ISC-GEM criterion for a usable depth (Di Giacomo et al., 2018).
- Number of defining phases: ≥ 6 minimum; ≥ 20 for well- constrained depth.
- RMS travel-time residual: < 1.0 s for well-located events (Bondár and Storchak, 2011); the ISC bulletin flags events with $RMS > 2.0$ s as poorly located (Storchak et al., 2020).
- Horizontal uncertainty: expressed as the semi-major axis of the 90 % confidence ellipse in km (Bondár et al., 2004).

These metrics should be used to rank competing solutions within a merged event group (section 1.9), not merely to flag events after the fact.

1.4.3 Depth fixing and its consequences

A significant fraction of agencies fix the focal depth at a standard value when insufficient depth-sensitive phases are available. Engdahl et al. (2020) estimate that approximately 30 % of events in the pre-1990 ISC bulletin have fixed depths. Fixed-depth solutions carry no depth uncertainty information and should be assigned lowest priority in any depth-selection step. The depth type (free or fixed, and if fixed, the value and reason) should be preserved as event metadata to enable downstream filtering (Di Giacomo et al., 2018; Engdahl et al., 2020).

The three most common fixed-depth conventions are 10 km (shallow crustal default, used by many regional agencies), 33 km (Conrad discontinuity, a historical NEIC convention), and 100 km (subduction-zone intermediate-depth default). Events with these three depths should be treated with heightened scrutiny in any depth-sensitive analysis (Engdahl et al., 1998).

1.5 The ISC-GEM Reference Methodology

The ISC-GEM Global Instrumental Earthquake Catalog (Storchak et al., 2013; Di Giacomo et al., 2018) is the most thoroughly documented end-to-end merging pipeline in the seismological literature. Its fifth revision (Di Giacomo et al., 2018) covers 1904–2014 for $M_w \geq 5.5$ globally and $M_w \geq 5.0$ over continental regions, with subsequent extensions to 2021. The pipeline is described in sufficient detail to serve as a reference implementation for any automated merging system.

1.5.1 Data sources and ingestion

Source data include the ISC bulletin (phase data and agency origin solutions from 1904 onward), the NEIC PDE and predecessors, GCMT solutions from 1976 (Dziewonski et al., 1981; Ekström et al., 2012), regional CMT and MT catalogues (GeoNet, INGV, GEOFON, ETH, NIED), and macroseismic intensity data for the pre-1960 period (Storchak et al., 2013). Each source is assigned a priority tier and a data quality flag based on known limitations (e.g. analogue vs. digital recording, reporting latency, known velocity model biases).

1.5.2 Event association

The ISC automated grouping algorithm compares origin times and epicentre coordinates from all contributing agencies and assembles candidate groups using a proximity-scoring system that weights azimuthal gap, phase count, distance range, and network geometry (Bondár and Storchak, 2011). Groups that appear to span two distinct earthquakes — identified by anomalously large within-group hypocentre scatter or by the same-agency rule — are split by the ISC analyst team (Storchak et al., 2020).

1.5.3 Relocation using all available phase data

The ISC does not simply select the best existing agency solution; it *re-locates every event* from scratch using all contributed phase data. The two-stage relocation procedure is:

1. **EHB algorithm** (Engdahl et al., 1998): an iterative relocation using the ak135 reference velocity model (Kennett et al., 1995) with improved phase identification, particularly for pP and sP depth phases. The EHB procedure reduced systematic depth biases in the pre-1990 bulletin by 10–30 km for subduction-zone events (Engdahl et al., 1998).
2. **ISC location algorithm** (Bondár and Storchak, 2011): fixes the depth output from EHB and further reduces systematic hypocentre bias by modelling the spatial correlation of travel-time prediction errors using an empirical covariance function estimated from the bulletin’s residual distribution.

The ISC-authored solution is designated *prime* when at least 6 defining phases are available and the azimuthal gap is $< 180^\circ$. Where these criteria are not met, the highest-ranked contributing agency solution is used as prime (Di Giacomo et al., 2018).

Crucially, *all* agency solutions are retained as secondary (non-prime) hypocentres in the ISC bulletin. This traceability model — derived from the practice of the original ISS — ensures that all downstream applications can reproduce the selection logic and compare solutions across agencies (Storchak et al., 2020).

1.5.4 Magnitude assembly in ISC-GEM

ISC-GEM reports a single M_w per event using the following protocol (Di Giacomo et al., 2018):

1. **Direct M_w** from global moment-tensor catalogs: GCMT (Dziewonski et al., 1981; Ekström et al., 2012) (preferred), USGS W-phase, EMSC MT, GeoNet CMT, ETH CMT.
2. **M_w proxy from M_s** : the ISC re-computes M_s using an *alpha-trimmed median* (10% trimming on each tail) of individual station amplitudes, suppressing outliers from poorly calibrated stations. The resulting M_s is converted to M_w via an exponential (non-linear) regression.
3. **M_w proxy from m_b** : similarly trimmed and converted via a non-linear regression that accounts for m_b saturation above $\approx M_w 6.2$.

The alpha-trimmed median is a robust location estimator in the statistical sense: its breakdown point (the fraction of outliers that can corrupt the estimate) is 10 %, compared with 0 % for the simple mean (Huber, 1981). For the station-level amplitude data used to compute M_s , where a single poorly calibrated station can deviate by 1–2 magnitude units, this robustness is practically important.

1.6 Magnitude Homogenisation: Theory and Practice

1.6.1 The physical basis for moment magnitude

Moment magnitude M_w was introduced by Hanks and Kanamori (1979) as a logarithmic measure of the seismic moment M_0 (measured in dyne-cm):

$$M_w = \frac{2}{3}(\log_{10} M_0 - 16.1). \quad (8)$$

The seismic moment $M_0 = \mu A \bar{D}$, where μ is the shear modulus, A is the fault rupture area, and $\bar{D}$ is the mean slip, is a physically meaningful measure of earthquake size that does not saturate at large magnitudes (Kanamori, 1977). In contrast, m_b saturates above $M_w \approx 6.0$ – 6.2 (because short-period body waves become insensitive to the growing fault area), and M_s saturates above $M_w \approx 8.0$ – 8.2 (Bormann, 2012).

The saturation problem directly affects merging: a network reporting $m_b = 6.5$ for an M_w 7.5 event is not recording a different earthquake; it is simply using a saturated scale. Magnitude-based filtering or association thresholds that do not account for scale-dependent saturation will therefore misclassify or incorrectly reject legitimate matches (Bormann, 2012; Di Giacomo et al., 2018).

1.6.2 Regression method: orthogonal vs. ordinary least squares

The fundamental flaw in using Ordinary Least Squares (OLS) for magnitude conversion is that OLS minimises the sum of squared vertical residuals, implicitly assuming the predictor variable is measured without error. When both the predictor M_x and the response M_w carry measurement error of comparable magnitude, OLS produces a biased slope estimator (Fuller, 1987):

$$\hat{\beta}_{\text{OLS}} = \beta \cdot \frac{\sigma_{M_x}^2}{\sigma_{M_x}^2 + \sigma_\varepsilon^2}, \quad (9)$$

where β is the true slope, $\sigma_{M_x}^2$ is the variance of the true predictor, and σ_ε^2 is the measurement error variance. The ratio is always ≤ 1 , so OLS systematically underestimates the true conversion slope — an *attenuation bias* that compresses the converted magnitude range and biases b -value estimates from the homogenised catalog (Grünthal et al., 2013; Carroll et al., 2006).

Generalised Orthogonal Regression. The solution is **Generalised Orthogonal Regression** (GOR), which minimises the sum of squared perpendicular distances weighted by the ratio $\eta = \sigma_{M_w}^2 / \sigma_{M_x}^2$ (Grünthal et al., 2013; Akkar et al., 2014; Cotton et al., 2006). The GOR slope is:

$$\hat{\beta}_{\text{GOR}} = \frac{(SYY - \eta SXX) + \sqrt{(SYY - \eta SXX)^2 + 4\eta SXY^2}}{2 SXY}, \quad (10)$$

where SXX , SYX , SXY are the corrected sums of squares and cross-products. When $\eta = 1$ (equal error variances), GOR reduces to orthogonal (major-axis) regression, and when η equals the ratio of the total variances SYX/SXX it reduces to the geometric mean (GM) regression. When $\eta \rightarrow 0$ (the response M_w is error-free), the slope tends to SYX/SXY , the reverse regression of M_x on M_w ; when $\eta \rightarrow \infty$ (the predictor M_x is error-free), it tends to SXY/SXX , ordinary least squares (Fuller, 1987).

The SHEEC (SHARE European Earthquake Catalogue) project (Grünthal et al., 2013) provides published GOR coefficients for M_L , M_s , and m_b to M_w conversion for European seismicity, estimated from the overlap period of the EMS98 national catalogs and the ISC-GEM dataset. These coefficients are recommended as defaults for European and Mediterranean merging applications (Grünthal et al., 2013).

Non-linear functional forms. For m_b and M_s , the ISC-GEM pipeline (Di Giacomo et al., 2018) chose exponential (non-linear) conversion models rather than linear GOR, on the basis that the physical saturation of both scales introduces a non-linear relationship at the extremes that linear regression cannot capture. Piecewise-linear models with optimised crossover points — as implemented in the CPTI15 Italian catalog pipeline using the Gasperini et al. (2013) framework — offer a practical middle ground that approximates non-linearity without the extrapolation risks of exponential models.

The Scordilis (2006) coefficients. The widely used Scordilis (2006) relations convert surface-wave and body-wave magnitudes, not local magnitudes: $M_w = 0.67 M_s + 2.07$ for $3.0 \leq M_s \leq 6.1$, $M_w = 0.99 M_s + 0.08$ for $6.2 \leq M_s \leq 8.2$, and $M_w = 0.85 m_b + 1.03$ for $3.5 \leq m_b \leq 6.2$. They were derived from a global dataset by orthogonal regression, which allows for error in both magnitudes, so the OLS attenuation bias of (9) does not apply to them. They are global averages: regional differences in magnitude calibration are not captured, and they should not be extrapolated outside their fitted ranges. No comparable global relation exists for M_L , whose calibration is network-specific; an $M_L \rightarrow M_w$ relation has to be derived regionally, for example from the overlap of a national catalog with moment-tensor solutions.

1.6.3 Magnitude type priority hierarchy

The ISC-GEM priority order (Di Giacomo et al., 2018), consistent with Grünthal et al. (2013) and the ISC bulletin operational practice, is:

1. M_w from GCMT waveform inversion (Dziewonski et al., 1981; Ekström et al., 2012) — gold standard; does not require conversion.
2. M_w from USGS/NEIC W-phase or broadband CMT solution.
3. M_w from regional CMT (GeoNet (GeoNet, 2025a), INGV, GEOFON/GFZ, ETH).
4. M_w proxy from M_s (preferred over m_b because M_s saturates at higher magnitude and the conversion scatter is smaller (Bormann, 2012)).

5. M_w proxy from m_b .
6. M_w proxy from M_L (most region-dependent; highly variable station corrections).
7. M_w proxy from M_d or M_c (duration / coda magnitude; least reliable, used only when nothing else is available).

Within each priority tier, the estimate with lower reported uncertainty is preferred (Di Giacomo et al., 2018). Where no uncertainty is reported, network authority from table 2 acts as a tie-breaker.

1.6.4 Bayesian magnitude merging across networks

Simple priority selection discards valid information when multiple authoritative M_w estimates exist for the same event. Si et al. (2024) proposed a Bayesian framework in which each network i 's observed magnitude is modelled as a linear function of the unknown true magnitude M_{true} :

$$M_{\text{obs},i} = a_i + b_i M_{\text{true}} + \varepsilon_i, \quad \varepsilon_i \sim \mathcal{N}(0, \sigma_i^2), \quad (11)$$

where a_i and b_i are network-specific intercept and slope (estimated from the overlap period), and σ_i^2 is the within-network reporting variance. A Gutenberg-Richter prior $p(M_{\text{true}}) \propto 10^{-bM_{\text{true}}}$ is placed on the true magnitude. The log-posterior, $-\frac{1}{2} \sum_i (M_{\text{obs},i} - a_i - b_i M_{\text{true}})^2 / \sigma_i^2 - b \ln(10) M_{\text{true}}$, is quadratic in M_{true} , so the posterior is Gaussian; completing the square gives the posterior mean

$$\hat{M}_{\text{true}} = \frac{\sum_i b_i (M_{\text{obs},i} - a_i) / \sigma_i^2 - b \ln(10)}{\sum_i b_i^2 / \sigma_i^2} = \sigma_{\text{post}}^2 \left[\sum_i \frac{b_i (M_{\text{obs},i} - a_i)}{\sigma_i^2} \right] - b \ln(10) \sigma_{\text{post}}^2, \quad (12)$$

where $\sigma_{\text{post}}^2 = (\sum_i b_i^2 / \sigma_i^2)^{-1}$ is the posterior variance: the Gutenberg-Richter prior shifts the likelihood estimate down by $b \ln(10) \sigma_{\text{post}}^2$. When two authoritative solutions agree within $0.1 M_w$ units, the posterior is insensitive to the choice of b ; when they disagree by $> 0.3 M_w$ units, a data quality flag should be raised regardless of which method is used (Si et al., 2024).

Uncertainty propagation for converted magnitudes. When M_w is estimated via conversion from M_x , the total uncertainty is (Grigoratos et al., 2023; Grünthal et al., 2013):

$$\sigma_{M_w}^2 = \sigma_{\hat{\alpha}}^2 + \left(\frac{\partial f}{\partial M_x} \right)^2 \sigma_{M_x}^2 + \sigma_{\text{model}}^2, \quad (13)$$

where $\hat{\alpha}$ is the conversion-equation intercept (with uncertainty $\sigma_{\hat{\alpha}}^2$ from the GOR fit), $f(M_x)$ is the conversion function, $\sigma_{M_x}^2$ is the individual measurement uncertainty of the predictor magnitude, and σ_{model}^2 is the residual scatter of the regression (model uncertainty). For the Scordilis $m_b \rightarrow M_w$ conversion with slope ≈ 0.85 , the conversion scales the input magnitude variance by $0.85^2 \approx 0.72$ (a reduction), but the model residual ($\sigma_{\text{model}} \approx 0.3$) typically dominates (Scordilis, 2006; Grünthal et al., 2013).

1.6.5 Combining macroseismic and instrumental estimates

Where both macroseismic ($M_{w,\text{macro}}$, derived from intensity data via Grünthal 1998) and instrumental ($M_{w,\text{inst}}$) estimates exist—a common situation for the historical period 1900–1960—the CPTI15 Italian pipeline (Rovida et al., 2016) uses inverse-variance weighting:

$$M_{w,\text{comb}} = \frac{M_{w,\text{macro}}/\sigma_{\text{macro}}^2 + M_{w,\text{inst}}/\sigma_{\text{inst}}^2}{1/\sigma_{\text{macro}}^2 + 1/\sigma_{\text{inst}}^2}. \quad (14)$$

This is the minimum-variance unbiased estimator under the assumption of independent Gaussian errors (Kay, 1993). For the Italian catalog, $\sigma_{\text{macro}} \approx 0.3\text{--}0.5$ and $\sigma_{\text{inst}} \approx 0.1\text{--}0.2$ (Rovida et al., 2016), so the combined estimate is dominated by the instrumental solution when both are available but gracefully falls back to the macroseismic estimate for events with no instrumental record.

1.7 Network Authority Hierarchies and Regional Overrides

1.7.1 The philosophical distinction: selection vs. re-analysis

Two fundamentally different philosophies underlie network authority models:

Selection model (ANSS ComCat). Each contributing network submits its preferred solution. A weight is assigned to each network per product type; the highest-weight solution is the “preferred” product (USGS, 2025). Where weights are equal, the most recently updated solution takes precedence. No re-analysis is performed by the aggregator.

Re-analysis model (ISC). The aggregator receives all raw phase data and performs its own relocation and magnitude computation using all contributed data simultaneously. The aggregator’s solution is prime; contributing solutions are retained as secondary (Storchak et al., 2020).

For research-grade applications, the re-analysis model is strongly preferred because it provides a single uniform solution computed with a common algorithm and velocity model, eliminating inter-agency biases (Storchak et al., 2013). For rapid operational applications where relocation latency is unacceptable, the selection model is appropriate.

1.7.2 Regional authority and the “authoritative network” principle

The ANSS ComCat implements the *authoritative network* principle: for earthquakes occurring within a defined geographic footprint, the regional specialist network is given the highest preferred weight, overriding the NEIC global solution (USGS, 2025). Examples include the Caltech/USGS Southern California Seismic Network (SCSN, code `ci`) for Southern California, the University of Utah Seismograph Stations (`uu`) for the Intermountain West, and GeoNet (`nz`) for Aotearoa New Zealand (GeoNet, 2025a).

The justification is that regional networks operate denser station arrays with local velocity models specifically calibrated to regional structure, producing lower-uncertainty locations and completeness to smaller magnitudes than any global

agency (Woessner and Wiemer, 2005; Hutton et al., 2010). For the New Zealand region, GeoNet (2025a) document that GeoNet solutions achieve completeness to approximately $M_L 1.5$ in the North Island, compared to the ISC’s practical completeness limit of $\approx M 4.5$ for the same region.

Table 2: Recommended global network authority hierarchy for parameter selection. Priority 1 = highest authority. Regional overrides (marked with †) take precedence over all entries above them in the table for earthquakes within the network’s geographic footprint.

Priority	Network / Agency	Code(s)	Notes
1	GCMT	gcmt	Waveform CMT; gold-standard M_w
2	ISC / ISC-GEM	isc	Re-located; full phase data retained
3	USGS NEIC	us	W-phase M_w ; global coverage
4	EMSC	emsc	Euro-Med aggregator
5	GEOFON / GFZ	gfz	Broadband; rapid solutions
6†	GeoNet / GNS	nz	Regional override: Aotearoa NZ
6†	JMA	jp	Regional override: Japan
6†	INGV	ingv	Regional override: Italy
6†	IGN	ign	Regional override: Spain
6†	BGR / GFZ	de	Regional override: Germany
7	IRIS / FDSN	ir	Waveform archive; no preferred solution

1.7.3 Temporal variation in network authority

Network capabilities change over time: networks expand or contract, recording equipment is upgraded, velocity models are revised. The authority hierarchy should therefore be applied with awareness of the recording era. For example, the GCMT catalog begins in 1976 (Dziewonski et al., 1981); for events before that date, the tier-1 authority falls to ISC or NEIC. GeoNet’s digital broadband network was not fully operational until the early 2000s; for pre-2000 New Zealand events, ISC or the GNS analogue bulletin should be used as the primary source (GeoNet, 2025a).

1.8 Regional Catalogue Merging: Case Studies

1.8.1 Dead Sea Transform region

Grigoratos et al. (2023) described an automated pipeline for the Dead Sea Transform region (DST-CAT), covering 1900–2017 and $M_w \geq 3.0$, assembled from ISC, the Geophysical Institute of Israel (GII), and several regional networks. Key methodological choices were:

- **Association:** 60 s / 100 km window with manual review of borderline cases.
- **Preferred solution:** specialist regional agency (GII) for events within the DST zone; ISC for the broader region; GCMT where available.
- **Magnitude conversion:** GOR-based M_L , M_s , and m_b to M_w relationships derived from the 1976–2017 overlap period, with uncertainties propagated via (13).
- **Era-dependent windows:** wider windows (90 s / 150 km) for the pre-1960 period to account for analogue timing errors.

The study explicitly noted the failure of fixed-window methods during the 1995 M_w 7.2 Aqaba mainshock sequence, where aftershock rates exceeded 10 events per hour for two days, and adopted a hybrid approach with manual review for the highest-rate periods.

1.8.2 Mexico unified catalogue

Sawires et al. (2019) assembled a unified Mexican seismicity catalog (1750–2018, $M \geq 3$) from seven source catalogs including the SSN (Servicio Sismológico Nacional), NEIC PDE, ISC, and the Reyes-Dávila regional bulletins. The deduplication procedure used a 60 s / 1° (≈ 111 km) window and a source priority list with the SSN as authoritative for events within Mexico. The study documents a magnitude change point in 1963 (transition from analogue to digital recording) requiring separate completeness analyses for the two periods.

1.8.3 Turkey homogeneous catalogue

Tan (2021) constructed a homogeneous M_w catalog for Turkey (1900–2018) from TURKNET, AFAD, ISC, GCMT, and several historical sources. The study used a 60 s / 100 km proximity window and applied GOR conversions with era-specific coefficients. The homogenised catalog was then declustered using the Reasenber (1985) algorithm before use in seismic hazard calculations.

1.8.4 Italy: CPTI15

The Parametric Catalogue of Italian Earthquakes (CPTI15, Rovida et al. 2016) is the most comprehensive regional merging effort for a seismically active European country, covering 1000–2014 with $M_w \geq 4.0$ for the instrumental period and extending to smaller magnitudes and earlier periods for historical events. It merges multiple national bulletins (INGV, GNDT, ING), the ISC bulletin, GCMT solutions, and a rich archive of macroseismic intensity data (Locati et al., 2016).

CPTI15 uses a piecewise-linear GOR framework (Gasperini et al., 2013) for magnitude conversion and the inverse-variance-weighted combination of (14) for events with both instrumental and macroseismic estimates. The catalogue is tied to the INGV ShakeMap intensity model to ensure consistency with ground-motion prediction equations used in Italian national hazard maps (Stucchi et al., 2013).

1.8.5 Aotearoa New Zealand

The GeoNet catalog (GeoNet, 2025a) is the authoritative source for New Zealand seismicity, providing real-time and reviewed solutions from the national seismograph network. Illsley-Kemp and Mestel (2025) recently produced a new high-precision catalog for the Taupō Volcanic Zone (TVZ) by re-processing GeoNet waveforms with a template-matching algorithm (Chamberlain et al., 2018), demonstrating that merged catalogs for New Zealand should treat the TVZ and broader crustal seismicity as distinct populations with different completeness thresholds.

The National Seismic Hazard Model (Gerstenberger et al., 2024) integrates the GeoNet catalog with paleo-seismic and geological data, requiring careful synchronisation of magnitude scales: GeoNet reports primarily M_L for small events and M_w for larger events, but the transition between scales is not uniformly applied across the catalog’s recording history (GeoNet, 2025b). For input to ground-motion models, all magnitudes must be expressed in M_w , necessitating the conversion procedures discussed in section 1.6.

1.9 Parameter Selection Within a Merged Group

Once event groups have been identified by one of the methods in section 1.3, parameter selection determines the single origin time, hypocentre, and magnitude that will represent the physical earthquake in the merged catalog.

1.9.1 Hypocenter and origin time

The preferred hypocentre is selected from the contributing solutions according to the priority order in table 2, subject to quality thresholds. The ISC-GEM criteria for accepting a prime hypocentre (Di Giacomo et al., 2018; Bondár and Storchak, 2011) are:

- Azimuthal gap $< 180^\circ$,
- At least 6 defining phases,
- Free depth (not fixed at a default value),
- RMS travel-time residual < 2.0 s.

The origin time is taken from the selected hypocentre *without averaging*. Averaging origin times across agencies is incorrect because agencies use different phase sets and velocity models; their origin times are not independent estimates of a common truth but are each internally consistent with their respective location procedures (Storchak et al., 2013).

Where no single agency meets all criteria, the highest-priority solution is used with a downgraded quality flag, and the criteria violation is recorded as event meta-data.

1.9.2 Depth selection

Among qualifying hypocentres, the solution with the smallest formal depth uncertainty is preferred. Where formal uncertainty is unavailable, the solution with the largest number of depth-sensitive phases (pP , sP , pwP , depth phases) is preferred. Depth solutions fixed at standard values (10, 33, 100 km) are given lowest priority and flagged with `depth_type = fixed` (Engdahl et al., 2020).

For events with no free-depth solution from any contributing agency, the depth should be reported as the fixed value of the best-available agency, with explicit documentation of the fixing convention used.

1.9.3 Magnitude selection and conversion

Apply the priority hierarchy of section 1.6.3. When multiple estimates exist at the same priority tier and formal uncertainties are available, the Bayesian posterior mean (12) is the optimal estimator (Si et al., 2024). The alpha-trimmed median (Di Giacomo et al., 2018) is an acceptable substitute when the Bayesian framework cannot be calibrated (e.g. for sparse regional catalogs without sufficient overlap for network intercept estimation). When neither is applicable, select the estimate from the highest-priority agency and document the choice.

1.9.4 Focal mechanism

For events where focal mechanisms are available from multiple sources, the following priority order is recommended based on the ISC-GEM protocol and established reliability assessments (Dziewonski et al., 1981; Ekström et al., 2012):

1. GCMT (waveform centroid moment tensor);
2. USGS/NEIC broadband CMT or W-phase solution;
3. GEOFON/GFZ regional moment tensor;
4. GeoNet CMT (for New Zealand events, GeoNet 2025a);
5. INGV CMT (for Italian events);
6. First-motion polarity solution with ≥ 20 station polarities;
7. Automated first-motion solution.

Within each tier, rank by: (i) variance reduction (higher is better), (ii) station polarity count (more is better), (iii) misfit (lower is better). All available mechanisms should be retained as supplementary products to enable downstream comparison studies (Storchak et al., 2020).

1.10 Declustering and Its Relationship to Merging

Declustering — the removal of aftershocks and foreshocks to obtain a Poissonian background catalogue — is a separate post-processing step from merging but is closely related in practice. Two points of interaction are important.

Merging before declustering. Declustering algorithms (e.g. Gardner and Knopoff 1974; Reasenberg 1985; Zaliapin and Ben-Zion 2013) require an internally consistent, homogeneous magnitude scale and a complete event listing to correctly identify cluster members. If two catalogs are merged *after* independent declustering, cluster members that appear in only one source catalog may be lost, biasing the background rate estimate (Hainzl et al., 2008). The correct order is: merge first (retaining all events from all sources), then decluster.

Duplicate detection vs. foreshock-aftershock discrimination. The nearest-neighbour proximity function (section 1.3.3) was introduced specifically to address the case where duplicate detection and aftershock identification are conflated (Vorobieva et al., 2022). This is a real operational hazard: aftershocks occurring within seconds of the mainshock are easily misidentified as duplicates by fixed-window methods during the most intense phase of a sequence. Systems that merge and decluster simultaneously (e.g. using a single proximity window) can lose genuine aftershocks and overcount the mainshock, distorting both the *b*-value and the Omori-law parameters of the sequence (Hainzl et al., 2008; Reasenberg, 1985).

1.11 Data Exchange Standards and Provenance

1.11.1 QuakeML and FDSN

The QuakeML 1.2 schema (ETH Zürich Seismology / QuakeML Working Group, 2015) is the de facto standard for representing earthquake event metadata in XML. It natively supports the representation of multiple competing origin solutions, multiple magnitude estimates per origin, focal mechanism solutions with full moment tensors, and phase arrival data. The `preferredOriginID` and `preferredMagnitudeID` attributes are the schema’s mechanism for expressing a selection decision, while retaining all contributing solutions as first-class objects.

FDSN event web services (U.S. Geological Survey, 2025) provide HTTP access to QuakeML records from participating agencies, enabling automated ingestion of source catalogs in a standard format. For merging systems that ingest from multiple FDSN nodes simultaneously, the `contributor` and `author` fields within each QuakeML origin provide the information needed to implement the same-agency rule (section 1.12.2) (ETH Zürich Seismology / QuakeML Working Group, 2015; U.S. Geological Survey, 2025).

1.11.2 Provenance tracking

Any merged catalog intended for research use should preserve the full derivation chain of each event record. The W3C PROV data model (World Wide Web Consortium, 2013) provides a formal language for expressing provenance graphs. For earthquake catalogs, provenance records should document:

- The source catalog identifier and version for each contributing record;
- The association algorithm and threshold values used;
- The selection rule applied for each parameter type (origin, magnitude, focal mechanism);

- The conversion equation and coefficients applied to any magnitude requiring homogenisation, with uncertainty estimates.

The FAIR data principles (Wilkinson et al., 2016) require that merged catalogs be assigned persistent identifiers (DOIs) with version tracking (DataCite, 2025). Where the input catalogs have their own DOIs, the provenance record should cite them explicitly to enable users to trace any merged event back to its source records (GO FAIR Initiative, 2025).

1.12 Conflict Detection and Quality Control

1.12.1 Within-group consistency flags

Quality flags should be computed for every merged event and stored as first-class metadata. Table 3 lists the recommended set, drawing on the ISC bulletin quality model (Bondár and Storchak, 2011; Storchak et al., 2020) and the ANSS ComCat practice (USGS, 2025). Flags are informational unless marked as “Error”; error-severity flags indicate that the group should be dissolved and re-evaluated before the event is included in the merged output.

Table 3: Recommended within-group quality flags for merged earthquake events. Severity: I = informational, W = warning, E = error (group should be dissolved).

Condition	Threshold	Severity	Reference
Azimuthal gap	$> 180^\circ$	W	Bondár and Storchak (2011)
Defining phase count	< 6	W	Bondár and Storchak (2011)
Depth fixed (not free)	any	I	Engdahl et al. (2020)
RMS travel-time residual	> 2.0 s	W	Storchak et al. (2020)
Horizontal uncertainty	> 50 km	W	Bondár et al. (2004)
$ \Delta M_w $: two authoritative sources	> 0.3	W	Si et al. (2024)
$ \Delta M_w $: converted vs. observed	> 0.5	W	Grünthal et al. (2013)
Magnitude range in group	> 1.0 (for $M < 5$)	W	Vorobieva et al. (2022)
Group size (events from different networks)	> 10	E	Patton et al. (2016)
Same network appears twice in group	any	E	Patton et al. (2016)
Spatial epicentre spread	> 100 km for $M < 5$	W	Patton et al. (2016)

(continued)

Condition	Threshold	Severity	Reference
Depth range within group	$> 50 \text{ km}$ (shallow)	W	Di Giacomo et al. (2018)
Wide time spread (informational)	$> 30 \text{ s}$ for $M < 5$	I	Di Giacomo et al. (2018)

1.12.2 The same-agency rule

If the same network contributes two records to a single merged group, the group almost certainly spans two distinct physical earthquakes. This is because any competent network applies intra-network deduplication upstream of its data submission; consequently, two records from the same network in the same time- distance window represent genuinely distinct detected events, not two reports of one event (Patton et al., 2016). The group should be dissolved and a secondary pairwise pass performed on its members, constrained by the same-agency rule as in the ILP constraint (7).

The platform implements this rule with one refinement. A single catalog never contributes two reports to a group, but two different catalogs may carry the same agency’s solution (two download vintages of one bulletin, say). Before any strategy runs, the reports of one agency in a group are compared on their agency event identifiers (the QuakeML event public identifier, else the source identifier), reduced to a bare form so that differently spelled copies of one identifier agree. A report counts as the agency’s only when its catalog is that agency’s and its own origin author, if stated, is the same agency; a bulletin row reproducing another agency’s hypocentre does not. Equal identifiers mark vintages of one solution: only the most recently determined takes part in selection, averaging and the validity checks, and the older vintages are kept in the event’s lineage, marked as superseded. Differing identifiers of the same kind are two earthquakes: the group fails validation and is dissolved and re-paired through the failed-group recovery below, as the literature recommends. Without comparable identifiers the reports are treated as independent, as reports of two agencies would be.

1.12.3 Failed-group recovery

When a group fails validation (e.g. because of a same-agency clash or anomalous magnitude range), the correct response is not to emit all members as singletons, but to attempt a constrained re-grouping. A one-level secondary greedy pass within the failed group — where the same-agency rule is enforced as a hard constraint — can recover valid sub-groups (Patton et al., 2016; Vorobieva et al., 2022). Sub-groups that still fail validation after the secondary pass are emitted as singletons with a QC flag recording the failed-group origin.

1.12.4 Magnitude completeness audit

The magnitude of completeness M_c should be estimated for the merged catalog in rolling temporal bins (5-year windows are recommended by Woessner and Wiemer 2005) using either the Maximum Curvature (MAXC) or Goodness-of-Fit (GFT) method (Wiemer and Wyss, 2000). Events below M_c must be flagged as potentially

incomplete in any bin where M_c is elevated and excluded from b -value or seismicity rate calculations that assume completeness (Aki, 1965). Sudden upward steps in M_c correspond to network deterioration or data gaps and should be documented alongside the network history that explains them (Woessner and Wiemer, 2005).

1.13 Algorithmic Recommendations

Table 4 maps the key algorithmic decisions to the methods recommended by the literature reviewed above, and summarises the current implementation status.

Table 4: Algorithmic decisions in earthquake catalog merging: literature recommendations, current implementation, and priority for future work. GOR = Generalised Orthogonal Regression; OLS = Ordinary Least Squares; DSU = Disjoint-Set Union.

Decision	Recommended	Current	Priority
Event association	Graph DSU (Patton et al., 2016)	Union-find over candidate pairs from different catalogs, joined closest first; one report per catalog per group; groups of three or more need a common centre; failed groups split	Done
Dense sequences	Nearest-neighbour η (Vorobieva et al., 2022)	Adaptive windows; closest-first pairing; preview flags ambiguous pairings	Low
Magnitude conversion	GOR / non-linear, regionally calibrated (Grünthal et al., 2013; Di Giacomo et al., 2018)	None for published magnitudes; global relations (Scordilis, 2006) and $M_L \approx M_w$ only in the consistency gate and to size the event for the type preference	High

(continued)

Decision	Recommended	Current	Priority
Multi-estimate averaging	α -trimmed median (Di Giacomo et al., 2018) or Bayesian (Si et al., 2024)	Magnitude selected by type; epicentre $1/\sigma^2$ -weighted (equal weights when a σ is missing), or the component-wise median epicentre as an alternative strategy	Medium
Authority table	User-editable, regional overrides (ISC/ComCat practice)	Editable in the application settings; NZ and Japan overrides by default; the table used is recorded per merged event	Done
Same-agency rule	Two reports of one network are two events (Patton et al., 2016)	Differing agency identifiers dissolve the group; identical identifiers mark the older vintage as superseded; otherwise independent	Done
Source provenance	All solutions retained (ISC model)	Every report kept per event, published one marked; exported per event	Done
Depth selection	Uncertainty-weighted (Engdahl et al., 2020)	Best-constrained free depth, with its own depth type and uncertainty; fixed depths only as a last resort	Low
Focal mechanism	GCMT > regional CMT	All mechanisms united; GCMT > USGS/NEIC > GEOFON > GeoNet > INGV > other MT > first motion (≥ 20 polarities) > automatic	Done

(continued)

Decision	Recommended	Current	Priority
QC flags	13-flag model (table 3)	9 flags implemented	Low
M_c audit	Rolling 5-yr bins (Wiemer and Wyss, 2000)	Not implemented (rate above M_c per bin for visual checks)	Medium
Declustering order	Merge first, then decluster	N/A (user responsibility)	Info
Provenance	W3C PROV records (World Wide Web Consortium, 2013)	Per-event lineage in every export; catalogue version and row checksum; no PROV records	Medium

The single highest-priority gap is magnitude homogenisation. The current implementation publishes each merged event’s magnitude unconverted; it uses the global relations of Scordilis (2006) for m_b and M_s , and $M_L \approx M_w$, only to judge whether the magnitudes in a duplicate group are consistent and how large the earthquake is, which sets the magnitude-type preference described below. A homogenised M_w catalog needs GOR-derived (Grünthal et al., 2013) or non-linear relations and, for the New Zealand region specifically, region-calibrated $M_L \rightarrow M_w$ coefficients estimated from the overlap of the GeoNet catalog with moment-tensor solutions, in preference to global coefficients (GeoNet, 2025b; Gerstenberger et al., 2024).

When it selects a magnitude for an averaged event, the current implementation also departs from the static ISC-GEM order of section 1.6.3, on saturation and calibration grounds, and publishes the chosen value as reported, not converted. The order depends on the size of the earthquake, taken as the median of its reported magnitudes expressed as M_w (M_L as M_w ; $M_w = 0.85 m_b + 1.03$; $M_w = 0.67 M_s + 2.07$ below M_s 6.2 and $0.99 M_s + 0.08$ from it; Scordilis, 2006). M_w always ranks first (the moment magnitudes M_{ww} , M_{wc} , M_{wb} and M_{wr} equally, and the proxies $M_w(m_B)$ and M_{wp} just below them). Below $M6.2$ the regionally calibrated M_L family (including GeoNet’s summary magnitude M) ranks ahead of m_b , m_B and m_{bLg} , then M_s , then M_d , because short-period m_b is poorly calibrated for small local events and M_s reads low there; from $M6.2$, where M_L approaches saturation and m_b has saturated, M_s ranks first, then m_B , M_L , m_b and m_{bLg} , and M_d . The switch sits at 6.2 because there the M_s relation becomes one to one while M_L still approximates M_w , so the published value stays within about 0.1 of M_w on either side of it. Ties go to the smaller uncertainty, then to the agency’s preferred magnitude, and rejected magnitudes are never selected. The ISC-GEM order ranks M_w *proxies* for conversion, which the implementation does not perform; a regionally calibrated conversion would make the distinction moot.

Event association has moved beyond the greedy forward sweep. Candidate pairs from different catalogs are joined with a union-find structure in increasing normalised space-time separation, so each report is claimed by its closest counterpart; unlike the plain connected-component approach of section 1.3.4, a group may not contain

two reports from one catalog (and a merge request that names one catalog twice is refused), and a group of three or more must have a member within the matching window of every other, so reports linked only through a chain of intermediate reports are not fused. Groups that fail the consistency checks are split into their largest self-consistent sub-groups, as recommended in the failed-group recovery above.

References

- Abe, K. (1981). Magnitudes of large shallow earthquakes from 1904 to 1980. *Physics of the Earth and Planetary Interiors*, 27(1):72–92.
- Aki, K. (1965). Maximum likelihood estimate of b in the formula $\log n = a - bm$ and its confidence limits. *Bulletin of the Earthquake Research Institute, University of Tokyo*, 43:237–239.
- Akkar, S., Sandıkkaya, M. A., and Bommer, J. J. (2014). Empirical ground-motion models for point- and extended-source crustal earthquake scenarios in Europe and the Middle East. *Bulletin of Earthquake Engineering*, 12(1):359–387.
- Bondár, I., Myers, S. C., Engdahl, E. R., and Bergman, E. A. (2004). Epicentre accuracy based on seismic network criteria. *Geophysical Journal International*, 156(3):483–496.
- Bondár, I. and Storchak, D. A. (2011). Improved location procedures at the International Seismological Centre. *Geophysical Journal International*, 186(3):1220–1244.
- Bormann, P., editor (2012). *IASPEI New Manual of Seismological Observatory Practice (NMSOP-2)*. Deutsches GeoForschungsZentrum GFZ, Potsdam.
- Carroll, R. J., Ruppert, D., Stefanski, L. A., and Crainiceanu, C. M. (2006). *Measurement Error in Nonlinear Models: A Modern Perspective*. Chapman & Hall/CRC, Boca Raton, FL, 2nd edition. Cited for treatment of attenuation bias in errors-in-variables regression.
- Chamberlain, C. J., Hopp, C. J., Boese, C. M., Warren-Smith, E., Chambers, D., Chu, S. X., Michailos, K., and Townend, J. (2018). EQcorrscan: Repeating and near-repeating earthquake detection and analysis in Python. *Seismological Research Letters*, 89(1):173–181.
- Cotton, F., Scherbaum, F., Bommer, J. J., and Bungum, H. (2006). Criteria for selecting and adjusting ground-motion models for specific target regions: Application to central Europe and rock sites. *Journal of Seismology*, 10(2):137–156.
- DataCite (2025). Versioning. <https://support.datacite.org/docs/versioning>. Guidance on DOI assignment and relations between dataset versions. Accessed: 2026-03-11.
- Di Giacomo, D., Engdahl, E. R., and Storchak, D. A. (2018). The ISC-GEM earthquake catalogue (1904–2014): Status after the extension project. *Earth System Science Data*, 10(4):1877–1899.

- Dziewonski, A. M., Chou, T. A., and Woodhouse, J. H. (1981). Determination of earthquake source parameters from waveform data for studies of global and regional seismicity. *Journal of Geophysical Research*, 86(B4):2825–2852.
- Ekström, G., Nettles, M., and Dziewonski, A. M. (2012). The global CMT project 2004–2010: Centroid-moment tensors for 13,017 earthquakes. *Physics of the Earth and Planetary Interiors*, 200–201:1–9.
- Engdahl, E. R., Di Giacomo, D., Sakarya, B., Gkarlaoui, C. G., Harris, J., and Storchak, D. A. (2020). ISC-EHB 1964–2016, an improved data set for studies of Earth structure and global seismicity. *Earth and Space Science*, 7(1):e2019EA000897.
- Engdahl, E. R., van der Hilst, R., and Buland, R. (1998). Global teleseismic earthquake relocation with improved travel times and procedures for depth determination. *Bulletin of the Seismological Society of America*, 88(3):722–743.
- EPOS (2025). Seismology — epos - the european plate observing system. <https://www.epos-eu.org/tcs/seismology>. Describes standardized, FAIR data services provided by the EPOS Seismology Thematic Core Service. Accessed: 2025-09-30.
- ETH Zürich Seismology / QuakeML Working Group (2015). Quakeml: A flexible xml representation of seismological data. <https://quake.ethz.ch/quakeml>. QuakeML overview and specification resources. Accessed: 2025-09-30.
- Fuller, W. A. (1987). *Measurement Error Models*. John Wiley & Sons, New York.
- Gardner, J. K. and Knopoff, L. (1974). Is the sequence of earthquakes in southern California, with aftershocks removed, Poissonian? *Bulletin of the Seismological Society of America*, 64(5):1363–1367.
- Gasperini, P., Lolli, B., and Vannucci, G. (2013). Empirical calibration of local magnitude data sets versus moment magnitude in Italy. *Bulletin of the Seismological Society of America*, 103(4):2227–2246.
- GeoNet (2025a). Earthquake catalogue. https://www.geonet.org.nz/data/types/eq_catalogue. Describes the GeoNet catalogue, Public ID policy, and access tools (Quake Search, WFS, FDSN, API). Accessed: 2025-09-30.
- GeoNet (2025b). Earthquake location information (seiscomp, locsat, nonlinloc, magnitudes). https://www.geonet.org.nz/data/supplementary/earthquake_location. Explains location algorithms, velocity models, and summary magnitude computation at GeoNet. Accessed: 2025-09-30.
- Gerstenberger, M. C. et al. (2024). The 2022 Aotearoa New Zealand national seismic hazard model: Process, overview, and results. *Bulletin of the Seismological Society of America*, 114(1):7–36.
- GO FAIR Initiative (2025). Fair principles. <https://www.go-fair.org/fair-principles/>. Accessed: 2026-03-11.
- Grigoratos, I., Poggi, V., Danciu, L., and Monteiro, R. (2023). Homogenizing instrumental earthquake catalogs – a case study around the Dead Sea Transform fault zone. *Seismica*, 2(2).

- Grünthal, G., editor (1998). *European Macroseismic Scale 1998 (EMS-98)*, volume 15. Cahiers du Centre Européen de Géodynamique et de Séismologie, Luxembourg.
- Grünthal, G., Wahlström, R., and Stromeyer, D. (2013). The SHARE european earthquake catalogue (SHEEC) for the time period 1900–2006 and its comparison to the European-Mediterranean earthquake catalogue (EMEC). *Journal of Seismology*, 17(4):1339–1344.
- Gutenberg, B. and Richter, C. F. (1954). *Seismicity of the Earth and Associated Phenomena*. Princeton University Press, Princeton, NJ, 2nd edition.
- Hainzl, S., Christophersen, A., and Enescu, B. (2008). Impact of earthquake rupture extensions on parameter estimations of point-process models. *Bulletin of the Seismological Society of America*, 98(4):2066–2072.
- Hanks, T. C. and Kanamori, H. (1979). A moment magnitude scale. *Journal of Geophysical Research*, 84(B5):2348–2350.
- Huber, P. J. (1981). *Robust Statistics*. John Wiley & Sons, New York.
- Hutton, K., Woessner, J., and Hauksson, E. (2010). Earthquake monitoring in Southern California for seventy-seven years (1932–2008). *Bulletin of the Seismological Society of America*, 100(2):423–446.
- Illsley-Kemp, F. and Mestel, E. (2025). A new consistent and high-precision earthquake catalogue for the Taupō Volcanic Zone, New Zealand. *Seismica*, 4(1).
- Kagan, Y. Y. (1991). Fractal dimension of brittle fracture. *Journal of Nonlinear Science*, 1(1):1–16.
- Kanamori, H. (1977). The energy release in great earthquakes. *Journal of Geophysical Research*, 82(20):2981–2987.
- Kay, S. M. (1993). *Fundamentals of Statistical Signal Processing: Estimation Theory*. Prentice Hall, Englewood Cliffs, NJ.
- Kennett, B. L. N., Engdahl, E. R., and Buland, R. (1995). Constraints on seismic velocities in the Earth from travel times. *Geophysical Journal International*, 122(1):108–124.
- Locati, M., Camassi, R., Rovida, A., Ercolani, E., Bernardini, F., Castelli, V., Caracciolo, C. H., Tertulliani, A., Rossi, A., Azzaro, R., and D’Amico, S. (2016). DBMI15, the 2015 version of the Italian Macroseismic Database. <https://emidius.mi.ingv.it/CPTI15-DBMI15/>. INGV. Accessed: 2026-03-29.
- Northern California Earthquake Data Center (2024). ANSS composite catalog search help. <https://www.ncedc.org/anss/catalog-search-help.html>. Documents the 16-second / 100-km duplicate-detection thresholds used in the ANSS Composite Catalog. Accessed: 2026-03-29.

- Obara, K., Kasahara, K., Hori, S., and Okada, Y. (2005). A densely distributed high-sensitivity seismograph network in Japan: Hi-net by National Research Institute for Earth Science and Disaster Prevention. *Review of Scientific Instruments*, 76(2):021301.
- Omori, F. (1894). On the aftershocks of earthquakes. *Journal of the College of Science, Imperial University of Tokyo*, 7:111–200.
- Pagani, M., Monelli, D., Weatherill, G., Danciu, L., Crowley, H., Silva, V., Henshaw, P., Butler, L., Nastasi, M., Panzeri, L., Simionato, M., and Vigano, D. (2014). OpenQuake Engine: An open hazard (and risk) software for the global earthquake model. *Seismological Research Letters*, 85(3):692–702.
- Patton, J. M., Guy, M. R., Benz, H. M., Buland, R. P., Erickson, B. K., and Kragness, D. S. (2016). Hydra—the National Earthquake Information Center’s 24/7 seismic monitoring, analysis, catalog production, quality analysis, and special studies tool suite. Open-File Report 2016-1128, U.S. Geological Survey.
- Reasenber, P. (1985). Second-order moment of central California seismicity, 1969–1982. *Journal of Geophysical Research: Solid Earth*, 90(B7):5479–5495.
- Rovida, A., Locati, M., Camassi, R., Lolli, B., and Gasperini, P. (2016). CPTI15, the 2015 version of the Parametric Catalogue of Italian Earthquakes. <https://emidius.mi.ingv.it/CPTI15-DBMI15/>. INGV. Accessed: 2026-03-29.
- Sawires, R., Santoyo, M. A., Peláez, J. A., and Corona Fernández, R. D. (2019). An updated and unified earthquake catalog from 1787 to 2018 for seismic hazard assessment studies in Mexico. *Scientific Data*, 6(1):241.
- Scordilis, E. M. (2006). Empirical global relations converting m_s and m_b to moment magnitude. *Journal of Seismology*, 10(2):225–236.
- Share, P.-E., Qiu, H., Vernon, F. L., Allam, A. A., Fialko, Y., and Ben-Zion, Y. (2022). General seismic architecture of the southern San Andreas fault zone around the Thousand Palms Oasis from a Large-N nodal array. *The Seismic Record*, 2(1):50–58.
- Si, Z., Zhuang, J., Gentili, S., Jiang, C., and Wang, W. (2024). A Bayesian merging of earthquake magnitudes determined by multiple seismic networks. *Seismological Research Letters*, 95(6):3221–3233.
- Storchak, D. A., Di Giacomo, D., Bondár, I., Engdahl, E. R., Harris, J., Lee, W. H. K., Villaseñor, A., and Bormann, P. (2013). Public release of the ISC-GEM global instrumental earthquake catalogue (1900–2009). *Seismological Research Letters*, 84(5):810–815.
- Storchak, D. A., Harris, J., Brown, L., Lieser, K., Shumba, B., and Di Giacomo, D. (2020). Rebuild of the Bulletin of the International Seismological Centre (ISC)—part 2: 1980–2010. *Geoscience Letters*, 7(1):18.
- Stucchi, M., Rovida, A., Gomez Capera, A. A., Alexandre, P., Camelbeeck, T., Demircioglu, M. B., Gasperini, P., Kouskouna, V., Musson, R. M. W., Radulian,

- M., Sesetyan, K., Vilanova, S., Baumont, D., Bungum, H., Fäh, D., Lenhardt, W., Makropoulos, K., Martinez Solares, J. M., Scotti, O., Živčić, M., Albin, P., Batllo, J., Papaioannou, C., Tatevossian, R., Locati, M., Meletti, C., Viganò, D., and Giardini, D. (2013). The SHARE european earthquake catalogue (SHEEC) 1000–1899. *Journal of Seismology*, 17(2):523–544.
- Tan, O. (2021). A homogeneous earthquake catalogue for Turkey. *Natural Hazards and Earth System Sciences*, 21(7):2059–2073.
- Tarjan, R. E. (1975). Efficiency of a good but not linear set union algorithm. *Journal of the ACM*, 22(2):215–225.
- U.S. Geological Survey (2025). Fdsn event web service specification — usgs implementation docs. <https://earthquake.usgs.gov/fdsnws/event/1/>. Reference for FDSN Event web service parameters and usage. Accessed: 2025-09-30.
- USGS (2025). ANSS comprehensive earthquake catalog (ComCat) documentation. <https://earthquake.usgs.gov/data/comcat/>. Accessed: 2025-10-21.
- Utsu, T. (1961). A statistical study on the occurrence of aftershocks. *Geophysical Magazine*, 30:521–605.
- Vorobieva, I. A., Gvishiani, A. D., Dzeboev, B. A., Dzeranov, B. V., Barykina, Y. V., and Antipova, A. O. (2022). Nearest neighbor method for discriminating aftershocks and duplicates when merging earthquake catalogs. *Frontiers in Earth Science*, 10:820277.
- Wiemer, S. and Wyss, M. (2000). Minimum magnitude of completeness in earthquake catalogs: Examples from Alaska, the western United States, and Japan. *Bulletin of the Seismological Society of America*, 90(4):859–869.
- Wilkinson, M. D., Dumontier, M., Aalbersberg, I. J., Appleton, G., Axton, M., Baak, A., Blomberg, N., Boiten, J.-W., da Silva Santos, L. B., Bourne, P. E., Bouwman, J., Brookes, A. J., Clark, T., Crosas, M., Dillo, I., Dumon, O., Edmunds, S., Evelo, C. T., Finkers, R., Gonzalez-Beltran, A., Gray, A. J. G., Groth, P., Goble, C., Grethe, J. S., Heringa, J., ’t Hoen, P. A. C., Hooft, R., Kuhn, T., Kok, R., Kok, J., Lusher, S. J., Martone, M. E., Mons, A., Packer, A. L., Persson, B., Rocca-Serra, P., Roos, M., van Schaik, R., Sansone, S.-A., Schultes, E., Sengstag, T., Slater, T., Strawn, G., Swertz, M. A., Thompson, M., van der Lei, J., van Mulligen, E., Velterop, J., Waagmeester, A., Wittenburg, P., Wolstencroft, K., Zhao, J., and Mons, B. (2016). The FAIR guiding principles for scientific data management and stewardship. *Scientific Data*, 3:160018.
- Woessner, J. and Wiemer, S. (2005). Assessing the quality of earthquake catalogues: Estimating the magnitude of completeness and its uncertainty. *Bulletin of the Seismological Society of America*, 95(2):684–698.
- World Wide Web Consortium (2013). Prov-o: The prov ontology. <https://www.w3.org/TR/prov-o/>. W3C Recommendation. Accessed: 2026-03-11.
- Zaliapin, I. and Ben-Zion, Y. (2013). Earthquake clusters in southern California I: Identification and stability. *Journal of Geophysical Research: Solid Earth*, 118(6):2847–2864.

Zaliapin, I., Gabrielov, A., Keilis-Borok, V., and Wong, H. (2008). Clustering analysis of seismicity and aftershock identification. *Physical Review Letters*, 101(1):018501.